General Guidelines

The following exercise is large, but can be done in under 4 hours with aggressive use of coding agents.

We encourage you to use AI **heavily** to implement this. If useful, we can give you an API key for any LLM of your choice with $50 of credits to use while implementing and testing this exercise.

Goal

Create a ChatGPT clone with advanced visualization capabilities.

Details

- The user should be able to ask questions about a data source they connect (can be excel files, mongodb, or **any datasource you choose** - you do not need to implement all of the above, rather implement any **single** datasource you choose)
- **When the user asks questions they should be able to get back text answers + visualizations - e.g. a pie chart that shows the answer**
- No need to implement for chat history
- No need to implement persistence
- No need to implement any feature in ChatGPT, other than what is described above
- Use Opus 4.6 as the LLM powering your ChatGPT clone

Grading

You will be graded based on:

- UX
- Product thinking
- Communication skills - this exercise was designed to test your performance on a real world team at Robusta, where you would communicate with a product manager and

customers on a regular basis. Therefore, if you have questions please ask them! Your ability to seek clarification when relevant is an important skill we are testing!

- Features you implement
- Architecture

Submission

Please submit your code as a github repo + share a SHORT Loom video:

1. Demoing what you built
2. Telling us how you built it (what coding agents did you use, how long did it take, etc)
3. Sharing any interesting architectural or product choices you made and why

Good luck!